

Ethical and Societal Implications of AI in Healthcare Diagnosis: A Case Study

Chinonyelum Igwe

ITAI 2372 – Artificial Intelligence Applications

Professor Ganti

July 3, 2026

Introduction

Artificial intelligence (AI) is changing the way doctors diagnose disease. Hospitals now generate huge amounts of data every day through electronic health records, medical imaging, lab tests, and even wearable devices, and it's becoming harder for physicians to keep up with all of it on their own. AI has stepped in as a tool that can sort through this data, spot patterns humans might miss, and help clinicians make faster, more informed decisions (Yu et al., 2018). It isn't meant to replace doctors so much as give them backup.

Healthcare diagnosis is one of the areas where AI has grown the fastest. Machine learning and deep learning models can now read medical images, patient histories, genetic data, and lab results to flag diseases that might otherwise slip through the cracks. Companies like Google Health, Aidoc, and Viz.ai have all built systems that assist physicians in detecting conditions such as breast cancer, strokes, and skin cancer more quickly and accurately (Esteva et al., 2017).

At the same time, using AI this way raises real ethical questions. Issues like patient privacy, bias, transparency, accountability, and public trust all come up once algorithms start playing a role in medical decisions. This case study looks at how AI diagnostic systems work, what benefits they bring, what ethical and societal concerns they raise, and what healthcare organizations can do to use this technology responsibly.

1. Overview of the AI System

AI in healthcare diagnosis refers to computer systems that use machine learning and deep learning to help physicians identify diseases and make clinical decisions. These systems can process far more medical information than a person could in the same amount of time, and

they're often able to catch subtle patterns that a physician might overlook during a routine review.

Traditionally, diagnosis has relied on doctors reading images, reviewing labs, and drawing on years of clinical experience. But the sheer volume of data modern healthcare produces makes it easy for something to get missed, and diagnostic errors are still a major problem worldwide, often leading to delayed treatment or unnecessary procedures. AI helps close that gap by processing large datasets quickly and picking out patterns that may point to disease (Yu et al., 2018).

Most of these systems are trained on thousands, sometimes millions, of medical records and images. During training, the algorithm learns to connect patient data with confirmed diagnoses, and once it's trained, it can evaluate new cases by comparing them to what it has already learned. Deep learning models, which are loosely modeled on the human brain, are especially good at reading complex images like X-rays, CT scans, MRIs, and pathology slides. Rajpurkar et al. (2017) built one such model, CheXNet, which could detect pneumonia from chest X-rays at a level comparable to a radiologist.

Real companies are already putting this to use. Google Health has developed models that can detect breast cancer in mammograms while cutting down on false positives and false negatives. Aidoc analyzes imaging in emergency departments to help radiologists catch urgent conditions like brain hemorrhages faster, and Viz.ai flags stroke patients from CT scans and alerts specialists almost immediately, which can make a real difference in how much damage a stroke ends up causing. Even with these advances, AI doesn't operate on its own, since physicians still review every recommendation and make the final call, because things like patient

communication and ethical judgment go beyond what a pattern-recognition system can do (Topol, 2019).

2. Benefits and Positive Impact

AI has the potential to make diagnosis faster, more accurate, and more accessible at a time when healthcare systems are dealing with growing patient populations and physician shortages. It's not there to replace doctors, but it can pick up on things that are easy for a person to miss.

The biggest benefit is probably improved diagnostic accuracy. AI models have shown strong results in detecting breast cancer, diabetic retinopathy, skin cancer, and heart conditions, sometimes catching early signs of disease before symptoms even show up (Esteva et al., 2017). Catching things earlier generally means better outcomes and higher survival rates.

AI also speeds things up. Radiologists look through hundreds of images a day, and having software that can analyze scans within seconds allows the most urgent cases to be prioritized. This matters most in emergency settings. Viz.ai, for instance, flags stroke symptoms on CT scans almost instantly so specialists can respond faster, and every minute saved lowers the risk of permanent damage.

There's a cost angle too. Delayed or wrong diagnoses lead to unnecessary tests, extra hospital visits, and longer recovery times, all of which add up. AI can cut down on some of that by improving precision earlier in the process, and hospitals benefit from not having physicians spend as much time on routine reviews.

AI is also helping close some access gaps. Rural areas and underserved communities often don't have enough specialists nearby, so AI tools paired with telemedicine let general practitioners offer a higher level of care even without a specialist in the room. On top of that, AI is speeding up medical research by helping identify disease patterns and supporting the shift toward more personalized medicine based on a patient's genetics and history (Yu et al., 2018).

3. Ethical Issues

As helpful as AI can be, it brings up some serious ethical questions, and healthcare is a field where getting things wrong actually costs people their health.

Algorithmic bias is probably the concern that comes up most. AI systems learn from historical data, and if that data doesn't represent a diverse patient population, the algorithm can end up performing worse for the groups it wasn't trained on well. That could mean missed or delayed diagnoses for certain populations, which only makes existing healthcare disparities worse (World Health Organization [WHO], 2021). Fixing this means using diverse datasets and regularly checking model performance across different groups.

Privacy is another big one. These systems need access to a lot of sensitive information, including health records, lab results, and genetic data, and any breach of that data could expose deeply personal information. Healthcare organizations have to follow regulations like HIPAA and use strong encryption, and patients deserve to know how their data is being used to train these systems (WHO, 2021).

Then there's the "black box" problem. A lot of deep learning models are extremely accurate but can't really explain how they reached a conclusion. In medicine, both doctors and patients usually want to understand the reasoning behind a diagnosis before making a decision

based on it, so limited explainability can chip away at trust. Researchers are working on explainable AI to try to fix this, but it's still an evolving area (Yu et al., 2018).

Accountability ties into all of this. If an AI system gets something wrong and a patient is harmed, it's not always clear who's responsible: the physician, the hospital, or the software company. Since AI is still a support tool and not an independent decision-maker, physicians need to stay in the loop and hold on to final responsibility (Topol, 2019). There's also the question of informed consent, since patients aren't always told when their data is being used to train these systems.

4. Societal Implications

The impact of AI in healthcare goes beyond the walls of a hospital. It touches jobs, trust, and how equally healthcare gets distributed.

A lot of people worry AI will eventually replace doctors and radiologists. Realistically, that's unlikely, since AI is good at recognizing patterns, but it can't replicate the judgment, empathy, and communication that come with treating a patient. What's more likely is that healthcare jobs shift rather than disappear, with professionals working alongside AI tools instead of being replaced by them, which means ongoing training will matter more going forward.

On the positive side, AI could help close healthcare gaps for rural and underserved communities by giving providers extra diagnostic support even when specialists aren't available nearby, and telemedicine paired with AI can extend care to areas that wouldn't otherwise have easy access to it.

That said, these benefits aren't guaranteed to be shared evenly. Wealthier hospitals have more money to invest in AI, while smaller or lower-income facilities may not be able to afford it, which risks widening the gap between well-resourced and under-resourced healthcare systems (National Institutes of Health [NIH], 2024). Public trust matters here too. Patients tend to be more comfortable with AI when they know a physician is still involved, and if people start to feel like decisions are being made by a computer with no human oversight, that could hurt confidence in the healthcare system as a whole.

5. Risks and Challenges

AI still has real limitations. It's only as good as the data it's trained on, so incomplete or biased training data can lead to inaccurate or inconsistent recommendations, which is dangerous if physicians start relying on it too heavily.

Cybersecurity is another major risk. Hospitals hold enormous amounts of sensitive patient data, which makes them a target for cyberattacks, and if that data is ever accessed or altered, it could compromise both patient privacy and patient safety. Strong encryption, regular security audits, and continuous monitoring are all necessary to keep that risk in check.

Overreliance is a related concern. AI should support a physician's judgment, not replace it, and doctors still need to weigh each patient's specific history and symptoms rather than taking an algorithm's output at face value.

Regulation is struggling to keep pace with the technology. The FDA has been working on guidelines for evaluating AI-based medical devices, but building rules that protect patients without slowing down innovation is a difficult balance to strike (U.S. Food and Drug Administration, 2024). AI models also need ongoing maintenance, since diseases and medical

knowledge change over time and a model that isn't regularly updated will slowly become less accurate.

6. Recommendations

A few practical steps can help healthcare organizations get the benefits of AI while limiting the risks.

Human oversight has to stay central. AI should support physicians' decisions, not replace them, and the final call on patient care should always come from a qualified professional. Developers also need to prioritize fairness by training models on diverse, representative datasets and testing regularly for bias across different patient groups, rather than assuming a model that works well for one population will work equally well for everyone.

Data privacy and cybersecurity protections need to be strong from the start: encrypted storage, restricted access, and full compliance with regulations like HIPAA, along with clear communication to patients about how their data is being used. Explainable AI should be a priority too, since physicians and patients are more likely to trust a system they can actually understand.

Regulators and governments need to keep building out clear guidelines for testing, validating, and monitoring AI systems in healthcare, and international cooperation will probably become more important as these tools spread globally (WHO, 2021). Finally, ongoing education for healthcare workers is essential. Physicians, nurses, and radiologists all need training on both what AI can do and where its limits are, so they can use it effectively instead of either ignoring it or over-trusting it.

Conclusion

AI has become one of the more significant advancements in modern healthcare, improving diagnostic accuracy, cutting down on wasted time and money, and expanding access to care in places that need it most. It isn't replacing physicians. If anything, it's giving them better tools to do their jobs.

But bringing AI into healthcare isn't just a technical challenge. It comes with real ethical and societal questions around fairness, privacy, transparency, and trust, and none of those get solved just by building a more accurate algorithm. Responsible use requires diverse data, human oversight, solid regulation, and ongoing evaluation to make sure these systems stay safe as they evolve.

Going forward, the goal shouldn't be AI replacing doctors, but AI and doctors working together, combining the speed and pattern recognition of the algorithm with the judgment and empathy that only a human provider can bring to patient care.

References

- Aidoc. (n.d.). AI-powered clinical solutions. <https://www.aidoc.com>
- Esteva, A., Kuprel, B., Novoa, R. A., Ko, J., Swetter, S. M., Blau, H. M., & Thrun, S. (2017). Dermatologist-level classification of skin cancer with deep neural networks. *Nature*, 542(7639), 115–118. <https://doi.org/10.1038/nature21056>
- Google Health. (n.d.). Artificial intelligence research in healthcare. <https://health.google>
- National Institutes of Health. (2024). Artificial intelligence in healthcare research. <https://www.nih.gov>
- Rajpurkar, P., Irvin, J., Zhu, K., Yang, B., Mehta, H., Duan, T., Ding, D., Bagul, A., Langlotz, C., Shpanskaya, K., Lungren, M. P., & Ng, A. Y. (2017). CheXNet: Radiologist-level pneumonia detection on chest X-rays with deep learning. *arXiv*. <https://arxiv.org/abs/1711.05225>
- Topol, E. (2019). *Deep medicine: How artificial intelligence can make healthcare human again*. Basic Books.
- U.S. Food and Drug Administration. (2024). Artificial intelligence and machine learning in medical devices. <https://www.fda.gov>
- Viz.ai. (n.d.). AI-powered stroke care. <https://www.viz.ai>
- World Health Organization. (2021). Ethics and governance of artificial intelligence for health. <https://www.who.int>
- Yu, K. H., Beam, A. L., & Kohane, I. S. (2018). Artificial intelligence in healthcare. *Nature Biomedical Engineering*, 2(10), 719–731. <https://doi.org/10.1038/s41551-018-0305-z>